

# Representation Models and Topic Modeling

Summer Methods Workshop

Taegyeon Kim | DHCSS, KAIST

Jul 31, 2025

# Agenda

## Things to be covered

- Static embeddings vs. contextual embeddings
- FNN
- Attention/transformer
- BERT
- BERTopic

# Static Embeddings

From contexts but not contextual

- Word2Vec, GloVe, Fast2Text
  - A word's embedding is derived from its context
- Different contexts do *not* lead to different embeddings  
→ *static* embeddings
- Even if a word has the same form, its meaning differs depending on the context

# Contextual Embeddings

## Example

- Embeddings for the same word differ by context
- Sentence 1: *Open a **bank** account*  
→  $e_{bank_{s1}} : [0.3, 0.9, \dots]$
- Sentence 2: *On the river **bank***  
→  $e_{bank_{s2}} : [0.8, 0.1, \dots]$



# Contextual Information in Neural Networks

Recent advances in NLP leverage neural networks to generate contextual representations of language

- Early models: RNNs, LSTMs, ULMFiT, ELMo
- Transformer-based models: BERT, GPT, T5, LLaMA, etc.
  - Use attention mechanisms to model complex contextual relationships
  - More effective and scalable

# Contextual Information in Neural Networks

Broadly, we can distinguish between two types of transformer-based models

- Representation models (e.g., BERT): trained to understand language by encoding input into rich, context-aware embeddings
  - Used primarily for producing representations that are useful for downstream tasks (e.g., topic modeling, classification)
- Generative models (e.g., GPT): trained to generate coherent text sequences by predicting the next word in context
  - These models learn representations as a means to an end (generation)
- This sessions focuses on BERT, one of the most influential representation models in NLP

# Attention and Transformer

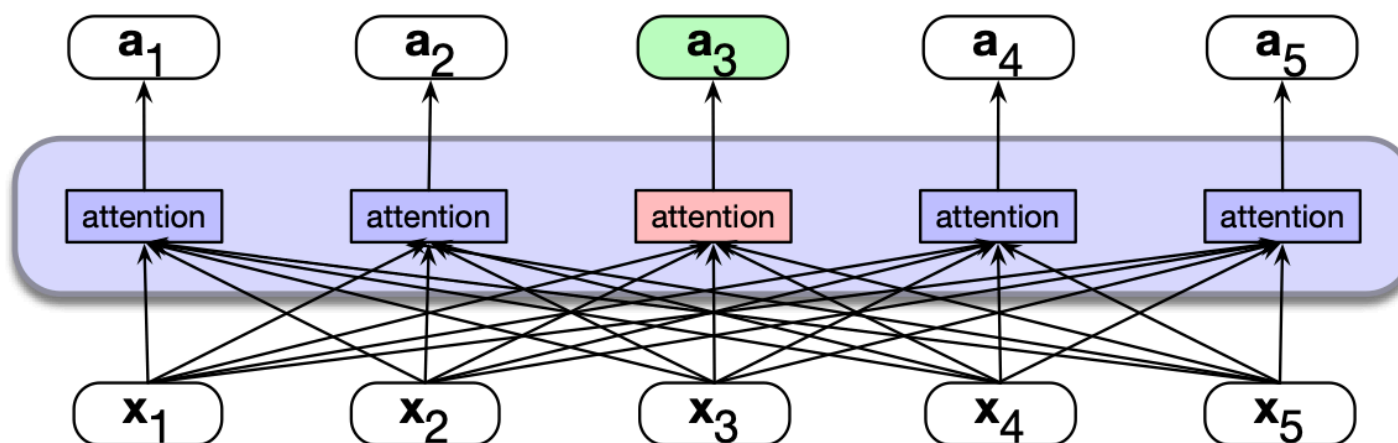
## Self-attention

- The primary goal of self-attention is to generate the representations of the tokens in a sequence
- Self-attention allows each word of the input sequence to ‘attend’ to (or reference) all other parts of the sequence (including self)
- In simple terms, attention is quantified as weights that indicate how much focus should be put on other tokens of the sequence in generating the representation of a given token
- E.g., “I took an umbrella from the old cabinet because it was raining”

# Attention and Transformer

## Self-attention

- Each token's representation is updated based on its relationship to all other tokens in the sequence
- Attention computes a weighted sum of the other token vectors, where the weights indicate relevance





# Attention and Transformer

## Self-attention

$$a_i = \sum_{j=1}^n \alpha_{ij} \cdot x_j$$

Where:

- $a_i$  is the updated representation for token  $i$
- $x_j$  is the input representation of token  $j$
- $\alpha_{ij}$  is the attention weight from token  $i$  to token  $j$

# Attention and Transformer

## Self-attention

- The weights are computed using similarity scores between tokens
- The dot product is used to measure how similar tokens  $i$  and  $j$  are

$$\text{score}(x_i, x_j) = x_i \cdot x_j$$

- Normalize these scores across all tokens using softmax
- This results in a probability distribution over all tokens  $j$ , indicating how much attention token  $i$  pays to each

# Attention and Transformer

## Self-attention

- This is a simplified illustration of self-attention
- In practice, transformer models (including generative models as well as representation models) use attention heads (Query, Key, and Value (QKV) projections)

# Attention and Transformer

## Another example

- *The chicken didn't cross the road because **it** was too tired.*
- *The chicken didn't cross the road because **it** was too wide.*

$e_{it_{s1}}$  and  $e_{it_{s2}}$  will have different representations after passing through attention layers

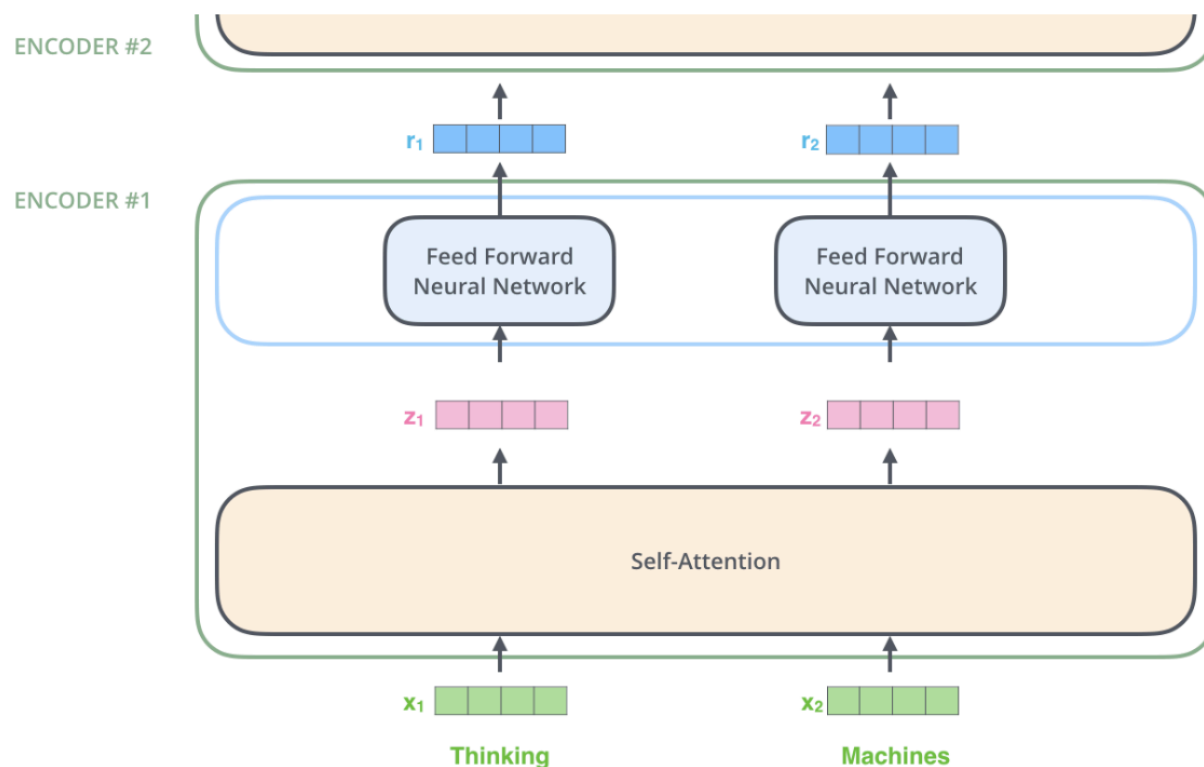
# Attention and Transformer

## What is Transformer

- Transformer is a type of neural network architecture with self-attention mechanisms
  - Introduced in 2017: [Vaswani et al. \(2017\)](#)
  - Cited 179,358 times as of May 12 2025
- Transformers consist of a stack of (encoder/decoder) layers main consisting of self-attention and feed forward neural networks

# Attention and Transformer

## Transformer encoder



Input

Embedding

Queries

Keys

Values

Score

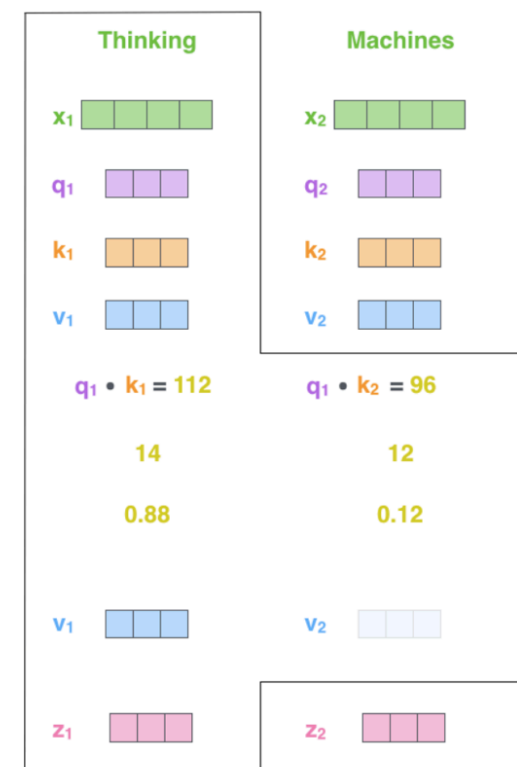
Divide by  $8 (\sqrt{d_k})$

Softmax

Softmax

X  
Value

Sum



# BERT

## Bi-directional Encoder Representations from Transformers

- A form of transformer trained as a language model
  - Two tasks: masked token prediction, next sentence prediction
  - Introduced in 2019: [Devlin et al. \(2018\)](#)
  - Cited 130,271 times as of May 12 2025
- Pre-trained on huge data sets ([BookCorpus](#) and Wikipedia)
- Two versions of the model introduced in the paper
  - BERT BASE (12 layers)
  - BERT LARGE (24 layers)

# BERT

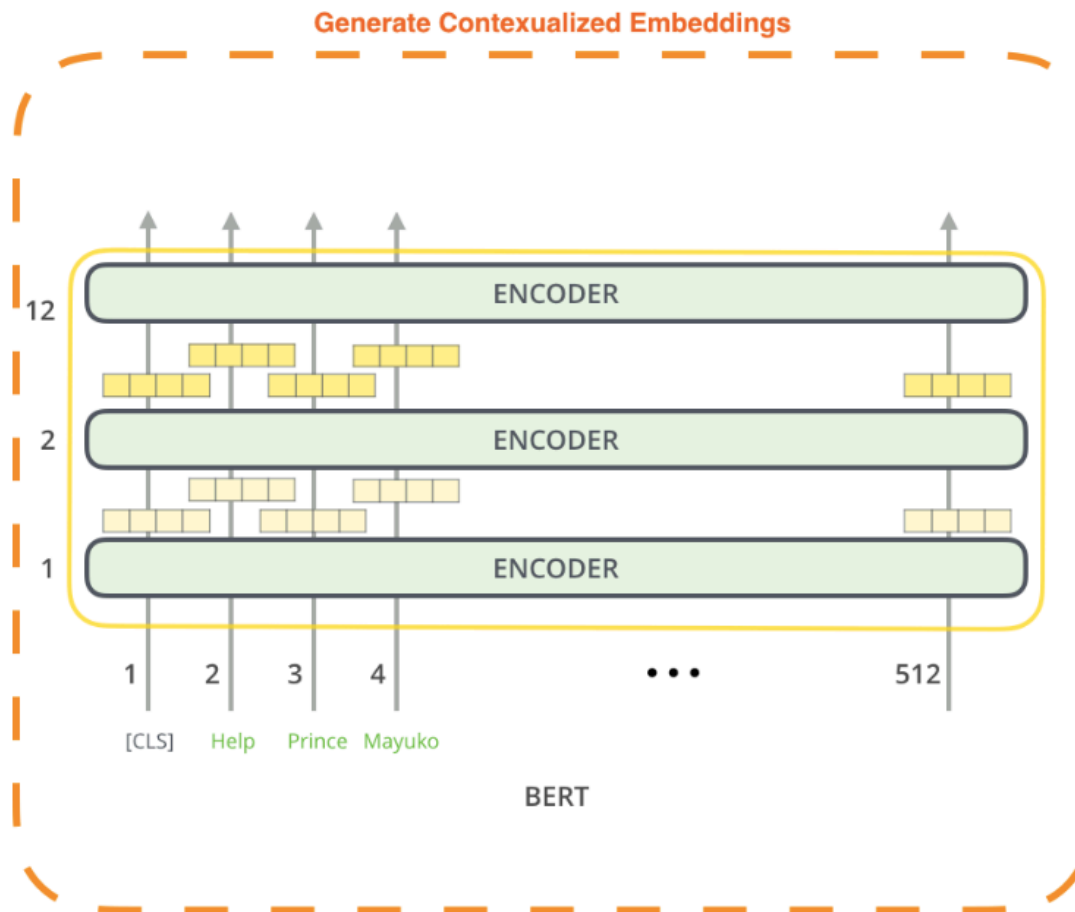
## Peak at under the hood

- BERT generates its own embeddings (from scratch) as part of its training process
- Stack of Transformer encoder layers involving “multi-head” self-attention
- Each layer passes its results through a feed-forward network, and then hands it off to the next encoder in the stack
- Each position outputs a vector of size 768 (BASE) or 1024 (LARGE)




# BERT

## Extracting embeddings



The output of each encoder layer along each token's path can be used as a feature representing that token.



But which one should we use?

# Topic Modeling with Contextual Embeddings

(Some of) practical usage in applied research

- Fine-tuning for text classification
  - E.g., identifying statements and stance on immigration  
(p. 7 *Classification* in [Card et al. 2022](#))
- The pre-trained model (and embeddings) is of interest
  - E.g., identifying dehumanizing metaphors against immigrants  
(p. 8 *Measuring Dehumanization* in [Card et al. 2022](#))
- Topic models
  - E.g., BERTopic

# Topic Modeling with Contextual Embeddings

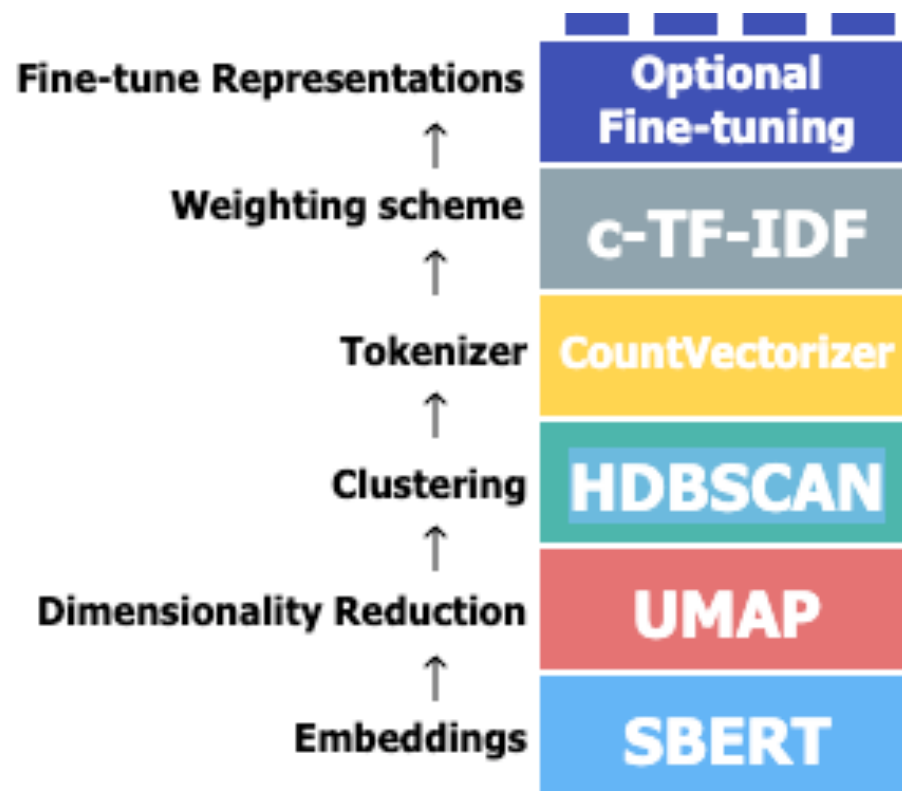
## Drawbacks for conventional models

- LDA, CTM, STM, DTM, etc.
- These are bag-of-words models
- Lack of capturing contextual information

# Topic Modeling with Contextual Embeddings

## Major steps in BERTopic

- Document embeddings: SBERT
- Dimensionality reduction: UMAP
- Clustering: HDBSCAN
- Topic representations: class-based TF-IDF



# Topic Modeling with Contextual Embeddings

## Characteristics of BERTopic

- Assumes no generative process
  - Unlike LDA, CTM, STM, etc.
- Semantically rich: can be good for short texts
- With embeddings from proprietary models, it is not quite transparent
- Guide coding involves detailed discussions of BERTopic in general and its modules, but *is modularity good?*

# Summary

- Modern NLP models based on transformer layers excel at capturing contextual information
  - Representations models are designed to understand and represent language
  - Generative models aim to generate coherent language
- Transformer layers powered by self-attention mechanisms form the core of these models' effectiveness
- Representation models are versatile
  - The embeddings themselves are valuable outputs
  - Useful for modeling the thematic structure of a corpus (i.e., topic modeling)
  - Applicable to downstream tasks like classification

# Guided Coding

- LDA basics: [BERTopic overview](#)
- LDA and BERTopic: [LDA/BERTopic on Stack Overflow data in Python](#)